

# BIT05 - Databanktechnologie

Jasper Anckaert

## Lecture 2 – Relational databases 2

# Previously

- Database
  - Collection of data that needs to be stored
  - Structured
  - Used everywhere
- Database system
  - Hardware – data – software – users
  - Storage space – quick – little redundancy – secure – clear structure
- Database Management System (DBMS)
  - Software application that interacts with the user, other applications, and the database itself to capture and analyse data
  - Storage – Retrieval – Manipulation – Authentication & authorization
- Relational databases – Relational Database Management System (RDBMS)
  - Enforce data integrity
  - Enforce referential integrity
  - Rules of E. Codd

# Previously

- MySQL
  - Install, connect to and secure server
  - User – host – database – table
  - Privileges
  - Options file
- Create database
- Grant privileges
- Show databases, tables columns, create statement

# Previously

- SQL
  - Data definition language
    - Statements to design database
    - CREATE, ALTER, DROP, ...
  - Data manipulation language
    - Statements to manage data
    - CRUD
    - SELECT, INSERT, UPDATE, DELETE
  - Data control language
    - Statements to manage database rights
    - GRANT, REVOKE

# Previously

## SQL: Structured Query Language

Diagram illustrating SQL clauses and their components:

- UPDATE clause { UPDATE movies Expression
- SET clause { SET rating = rating + 1
- WHERE clause { WHERE name = 'USA';

Annotations:

- Expression (under rating + 1)
- Expression (under 'USA')
- Predicate (under WHERE name = 'USA';)
- Statement (grouping all clauses)

# Previously

## Column types

- INT
  - Integer
  - SIGNED: -2 147 483 648 tot 2 147 483 647
  - UNSIGNED: 0 tot 4 294 967 295
  - TINYINT, BIGINT, SMALLINT
- FLOAT & DOUBLE
  - Numbers with decimal point
  - FLOAT: 7 digits after decimal point, DOUBLE: 15 digits after decimal point
- DATE
  - YYYY-MM-DD
  - DATETIME
    - YYYY-MM-DD HH:MM:SS
    - ! TIMESTAMP ! No dates < 1970 and > 2038

# Previously

## Column types

- VARCHAR & CHAR
  - String with a certain number of characters
  - Define max number of characters e.g. VARCHAR(200)
  - VARCHAR: up to 65 535 characters
  - CHAR: up to 255 characters, spaces are added to reach required length

CHAR(10)



VARCHAR(10)



- VARCHAR is more efficient in storage, CHAR is faster for reading data
- Similar for INT vs BIGINT vs ...



# Previously

## Column types

- TEXT & BLOB
  - Used for texts that are not queried often or do not have to be searchable
  - BLOB for binary data (images, ...)
- ENUM
  - List of permitted values
    - E.g. Set of colours: 'red', 'green', 'blue'
  - Very efficient

# Previously

## Constraints

On top of column types, there are some additional requirements per column

- Primary key
  - Only 1 PK per table, all values must be unique
- UNIQUE
  - All values (or combinations) must be unique
- NOT NULL
  - Field can not be empty when adding data (empty = null)
- Default
  - Default value for a field
- Foreign key
  - Same constraints as referenced column
  - Security when adjusting linked data possible

# Previously

- INSERT

```
INSERT INTO tbl (col1, col2) VALUES (val1, val2);
```

- SELECT

```
SELECT columns FROM tbl;
```

- ORDER BY

```
SELECT columns FROM tbl ORDER BY col1 [asc|desc] [, col2 [asc|desc]...];
```

- Calculated rows

- Built in functions for numbers, strings, dates

- Column aliases

- Can be used in the ORDER BY clause

- WHERE

```
SELECT columns FROM tbl WHERE condition(s) [ORDER BY sortcol];
```

- NULL values

```
SELECT ... WHERE col IS [NOT] NULL;
```

```
SELECT ifnull(col, value) ...
```

# Previously

- AND, OR, NOT, XOR

- Boolean logic

- DISTINCT

`SELECT DISTINCT(cols) FROM ...`

- LIMIT, OFFSET

`SELECT ... LIMIT n [OFFSET x];`

- Aggregation

- Built in functions e.g. `count()`, `sum()`, `min()`, `max()`, ...

- GROUP BY

`SELECT [col,] aggregatefunctions FROM src [WHERE cond] GROUP BY col [ORDER BY ...];`

- HAVING

`SELECT [col, ] aggregatefunctions FROM src [WHERE cond1] GROUP BY col HAVING cond2 [ORDER BY ...];`

# Previously

## Execution order

1. Input columns are determined
2. WHERE – input columns are filtered
3. GROUP BY – sorting & grouping of filtered input
4. Aggregation functions are calculated
5. HAVING – aggregation functions are filtered
6. ORDER BY – output is sorted
7. LIMIT/OFFSET – output is chopped

# Relational databases with MySQL

## Exercises

- Return a list of the 3 biotypes with the most genes. Which biotype is a close fourth?
- Return a list of the number of genes per status
- Combining the 2 previous results, which biotype is most known?
- Select only those biotypes that cover at least 3% of the human genome (hint: 'size' in previous exercise, human genome is approx. 3 billion bp) and return this percentage.

# Relational databases with MySQL

## Database upgrade

- Download 2.sql
- Create a new database bioinf

```
mysql> CREATE database bioinf;
```
- Create the tables and insert the data

```
$ mysql bioinf < 2.sql
```

# Relational databases with MySQL

## Joins

- Relation databases model entities and their relationships
- Different entities: different tables
- Allow you to combine information across different tables



# Relational databases with MySQL

## Joins

- What if we want to expand our database with a trajectory and course info

Student_number	Name	Last_name	Birthdate	Trajectory	Course
0293826	John	Doe	1991-10-02	FBT	Databases
0293749	Mel	Trotter	1991-04-11	MLT	Databases
0328273	Bill	Schuette	1990-12-01	MLT	Databases
0293826	John	Doe	1991-10-02	FBT	Scripting
0293826	John	Doe	1991-10-02	FBT	Linux

- Problem: redundant information
  - Waste of space
  - Error prone

# Relational databases with MySQL

## Joins

- Solution: relational databases with a **foreign key**

*Students*

Student_number	Name	Last_name	Birthdate	Trajectory_ID	Course
0293826	John	Doe	1991-10-02	1	Databases
0293749	Mel	Trotter	1991-04-11	2	Databases
0328273	Bill	Schuette	1990-12-01	2	Databases
0293826	John	Doe	1991-10-02	1	Scripting
0293826	John	Doe	1991-10-02	1	Linux

- Important: Data types of linked columns have to be equal! Foreign key is typically primary key of other table

*Trajectories*

ID	Trajectory
1	FBT
2	MLT

# Relational databases with MySQL

## Joins

*Students*

Student_number	Name	Last_name	Birthdate	Trajectory_ID
0293826	John	Doe	1991-10-02	1
0293749	Mel	Trotter	1991-04-11	2
0328273	Bill	Schuette	1990-12-01	2

*Course\_of\_student*

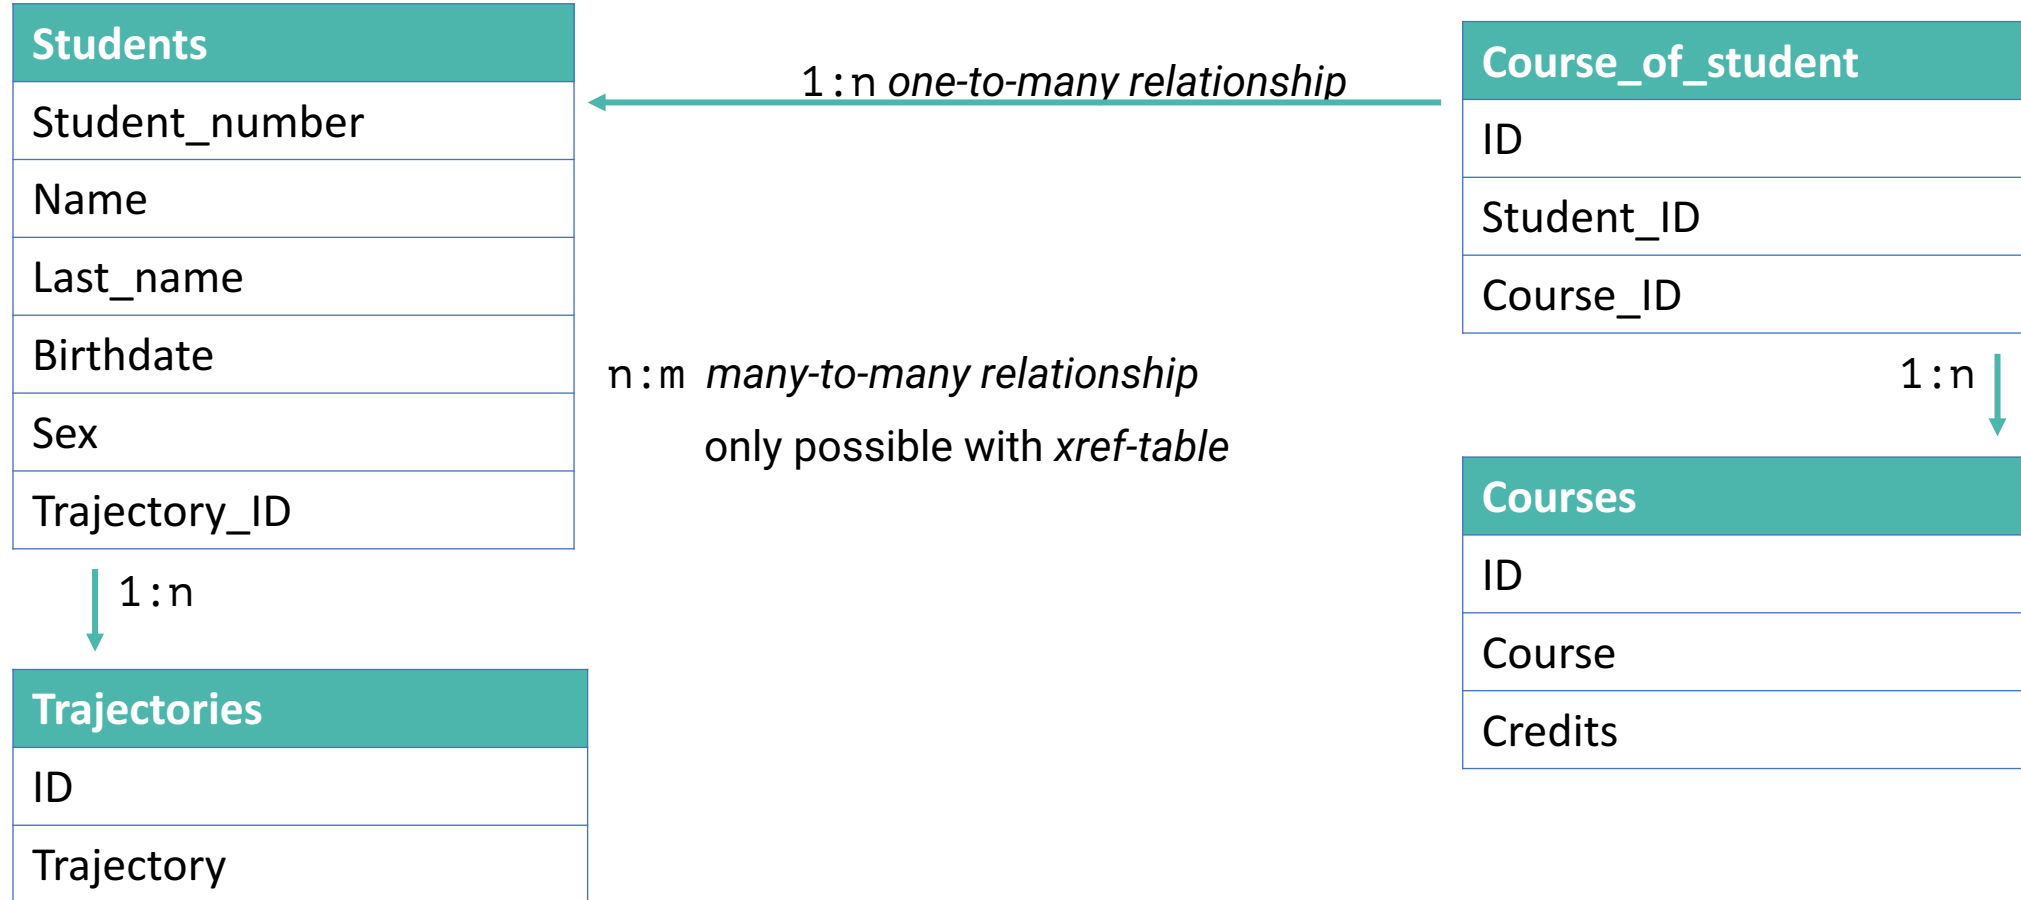
Student_id	Course_ID
0293826	1
0293826	2
0293826	3
0293749	1
0328273	1

*Courses*

ID	Course
1	Databases
2	Scripting
3	Linux

# Relational databases with MySQL

## Joins



Database model= Entity Relationship Diagram (ERD)

# Relational databases with MySQL

Joins – Information is spread across multiple tables

- ~~Create script that retrieves tables separately~~
- JOIN tables with query

Students
Student_number
Name
Last_name
Birthdate
Sex
Trajectory_ID
Trajectories
ID
Trajectory



Student_number	Name	Last_name	Birthdate	Course
0293826	John	Doe	1991-10-02	Databases
0293749	Mel	Trotter	1991-04-11	Databases
0328273	Bill	Schuette	1990-12-01	Databases
0293826	John	Doe	1991-10-02	Scripting
0293826	John	Doe	1991-10-02	Linux

Data representation does not have to be equal to the way the data is stored

# Relational databases with MySQL

## Joins

- Retrieve linked rows from different tables with JOIN

```
mysql> SELECT * FROM Students  
      JOIN Trajectories ON Students.Trajectory_ID = Trajectories.ID;
```

Students

Student_number	Name	Last_name	Birthdate	Trajectory_ID
0293826	John	Doe	1991-10-02	1
0293749	Mel	Trotter	1991-04-11	2
0328273	Bill	Schuette	1990-12-01	2
0324312	Penelope	Tracy	1989-07-24	<i>NULL</i>

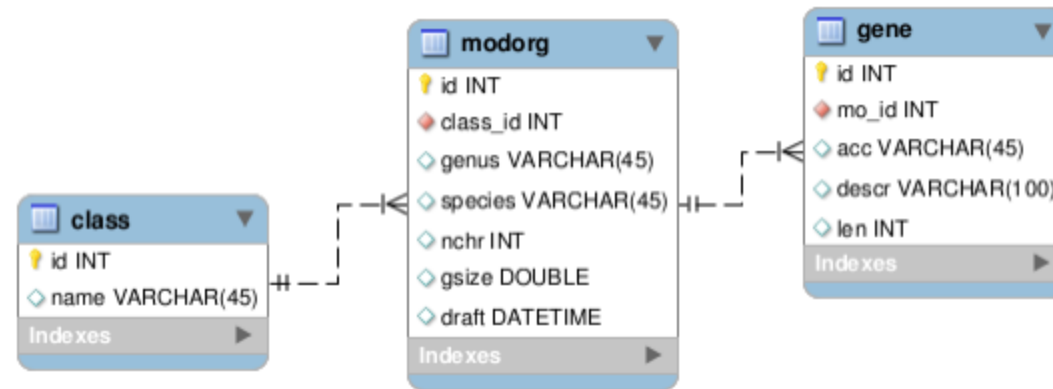
Trajectories

ID	Trajectory
1	FBT
2	MLT
3	TI

# Relational databases with MySQL

Joins – back to our database

- `modorg.class_id` is a *foreign key* that references `class.id`
- `gene.mo_id` is a *foreign key* that references `modorg.id`



# Relational databases with MySQL

## Joins – Cartesian product

- Multiple tables in a query → database server generates all possible combinations  
= Cartesian product

table `class` has 6 rows

table `modorg` has 10 rows

query `SELECT * FROM modorg, class` has 60 rows



# Relational databases with MySQL

## Joins

- Filtered Cartesian product

```
SELECT * FROM modorg, class WHERE modorg.class_id = class.id;
```

```
SELECT * FROM modorg [INNER] JOIN class ON modorg.class_id =  
class.id;
```

# Relational databases with MySQL

## Joins

- Avoid ambiguity

```
mysql> SELECT id, name, genus, species FROM modorg, class WHERE  
modorg.class_id = class.id;
```

```
ERROR 1052 (23000): Column 'id' in field list is ambiguous
```

- Ambiguous columns must be qualified. And additionally you can choose an alias:

```
mysql> SELECT modorg.id as mo_id, name, genus, species FROM  
modorg, class WHERE modorg.class_id = class.id;
```

# Relational databases with MySQL

## Joins – data source alias

- Join table with itself or select data from subquery
  - Use alias for data source

```
mysql> SELECT a.col, b.col FROM src1 [as] a, src2 [as] b WHERE ...
```

# Relational databases with MySQL

## Joins – data source alias

- For each class, give the class name, organism name and date of the organism that was sequenced first

```
mysql> SELECT class_id, min(draft) as dr FROM modorg GROUP BY  
class_id;
```

- add class name: join with table `class`
- add organism name: join with table `modorg`

# Relational databases with MySQL

## Joins – data source alias

- Add class name

```
mysql> SELECT name, dr FROM  
      (SELECT class_id, min(draft) as dr FROM modorg GROUP BY  
       class_id) as s, class WHERE s.class_id=class_id;
```

- add organism name: join with table modorg

```
mysql> SELECT name, concat(genus," ",species) as org_name, dr  
      FROM (SELECT class_id, min(draft) as dr FROM modorg GROUP BY  
            class_id) as s, class, modorg WHERE s.class_id=class.id AND  
            s.dr=draft;
```

# Relational databases with MySQL

## Exercises

- For all rows in table gene, show
  - Organism name
  - Class name
  - Accession
  - Length
  - Description of the gene

# Relational databases with MySQL

## Joins – 4 types

- INNER JOIN
  - Only rows present in both tables
- LEFT JOIN
  - All rows from left table, even without linked data in right table
- RIGHT JOIN
  - All rows from right table, even without linked data in left table
- OUTER JOIN
  - All rows from both tables
  - Doesn't exist in MySQL

# Relational databases with MySQL

## Joins – INNER JOIN

Students

Student_number	Name	Last_name	Birthdate	Trajectory_ID
0293826	John	Doe	1991-10-02	1
0293749	Mel	Trotter	1991-04-11	2
0328273	Bill	Schuette	1990-12-01	2
<del>0324312</del>	<del>Penelope</del>	<del>Tracy</del>	<del>1989-07-24</del>	<del>NULL</del>

Trajectories

ID	Trajectory
1	FBT
2	MLT
<del>3</del>	<del>TI</del>



Student_number	Name	Last_name	Birthdate	Trajectory
0293826	John	Doe	1991-10-02	FBT
0293749	Mel	Trotter	1991-04-11	MLT
0328273	Bill	Schuette	1990-12-01	MLT



# Relational databases with MySQL

## Joins – LEFT JOIN

Students

Student_number	Name	Last_name	Birthdate	Trajectory_ID
0293826	John	Doe	1991-10-02	1
0293749	Mel	Trotter	1991-04-11	2
0328273	Bill	Schuette	1990-12-01	2
0324312	Penelope	Tracy	1989-07-24	NULL

Trajectories

ID	Trajectory
1	FBT
2	MLT
3	TI



Student_number	Name	Last_name	Birthdate	Trajectory
0293826	John	Doe	1991-10-02	FBT
0293749	Mel	Trotter	1991-04-11	MLT
0328273	Bill	Schuette	1990-12-01	MLT
0324312	Penelope	Tracy	1989-07-24	NULL

# Relational databases with MySQL

## Joins – RIGHT JOIN

Students

Student_number	Name	Last_name	Birthdate	Trajectory_ID
0293826	John	Doe	1991-10-02	1
0293749	Mel	Trotter	1991-04-11	2
0328273	Bill	Schuette	1990-12-01	2
<del>0324312</del>	<del>Penelope</del>	<del>Tracy</del>	<del>1989-07-24</del>	<del>NULL</del>

Trajectories

ID	Trajectory
1	FBT
2	MLT
3	TI



Student_number	Name	Last_name	Birthdate	Trajectory
0293826	John	Doe	1991-10-02	FBT
0293749	Mel	Trotter	1991-04-11	MLT
0328273	Bill	Schuette	1990-12-01	MLT
NULL	NULL	NULL	NULL	TI

# Relational databases with MySQL

## Joins – OUTER JOIN

Students

Student_number	Name	Last_name	Birthdate	Trajectory_ID
0293826	John	Doe	1991-10-02	1
0293749	Mel	Trotter	1991-04-11	2
0328273	Bill	Schuette	1990-12-01	2
0324312	Penelope	Tracy	1989-07-24	<i>NULL</i>

Trajectories

ID	Trajectory
1	FBT
2	MLT
3	TI



Student_number	Name	Last_name	Birthdate	Trajectory
0293826	John	Doe	1991-10-02	FBT
0293749	Mel	Trotter	1991-04-11	MLT
0328273	Bill	Schuette	1990-12-01	MLT
0324312	Penelope	Tracy	1989-07-24	<i>NULL</i>
<i>NULL</i>	<i>NULL</i>	<i>NULL</i>	<i>NULL</i>	TI

# Relational databases with MySQL

## Joins – characteristics

- \* is used to select all columns from all tables
- Use *tbl.col* to specify column
- Very inefficient
  - A lot of memory, a lot of time
- Use ON to specify linked columns

# Relational databases with MySQL

## Execution order

1. Input columns are determined
  - a. JOIN clause
2. WHERE – input columns are filtered
3. GROUP BY – sorting & grouping of filtered input
4. Aggregation functions are calculated
5. HAVING – aggregation functions are filtered
6. ORDER BY – output is sorted
7. LIMIT/OFFSET – output is chopped

# Relational databases with MySQL

## Exercices

- Switch back to the `bioinf_testdb` database
  - Experiment with the 3 types of JOIN on your own tables
  - Examine the `gene`, `transcript` and `exon` table, how are they connected?
    - How many transcript does the *MALAT1* gene have?
    - Return the position of the exons of transcript 237999
  - Return the transcripts of the *TP53* gene
  - Return their exons as well
  - Find the longest spliced transcript of *TP53* (taking into account the intron-exon structure)
  - How many exons does each transcript have?

# Relational databases with MySQL

## Exercises

- What is the name of the gene associated with transcript 260392?
- A mutation was found on chromosome 20, position 44002590. Which gene(s) overlap(s) with this position?
  - Select only those genes that have exons that overlap with this mutation. Which genes are they?
- A biotype column can be found in both the gene and transcript table. Are there transcripts that have a different biotype from the gene they're part of? What are their names?
  - Does the same go for status?
- Which chromosome has the most genes and how many are there?
- Which exon is the largest in the genome and how many base pairs are there?
- Which transcript has the most exons and how many are there?

# Relational databases with MySQL

## Efficiency and speed

- Complex queries tend to become
  - Large system load
  - Slow interface, user has to wait
- Mind your choices of column types and included columns in the query
- Frequently used queries can be sped up
  - Views
  - Indices
  - Allow redundancy



# Relational databases with MySQL

## Views

- Re-use same query
- Query can be saved as a special table-like object
- Usually read-only data source
- Speed gain depends on use case
- MySQL
  - Virtual table
  - Used to serve up data in an orderly fashion
- Oracle i.a.
  - Materialized view: result of a SELECT query is stored
    - Very efficient

# Relational databases with MySQL

## Views

- Create a new view

```
mysql> CREATE VIEW viewname as SELECT ...
```

- Use a view as a table

```
mysql> SELECT ... FROM viewname WHERE ... ORDER BY ...
```

# Relational databases with MySQL

## Exercises

- Create a view (genevw) from the query in the previous exercise
- Select all genes containing hemoglobin in the description and sort the result set by gene length.

Next:

- What is the minimum and maximum gene length?
  - What is the average gene length?
  - And the standard deviation?
- Does the view show up when using  
`mysql> show tables;`

# Relational databases with MySQL

## Index

- Quickly find certain rows
- Stored separately
- Golden rule: Use indices on columns you will use in your WHERE clause
- Only 1 index per query is used

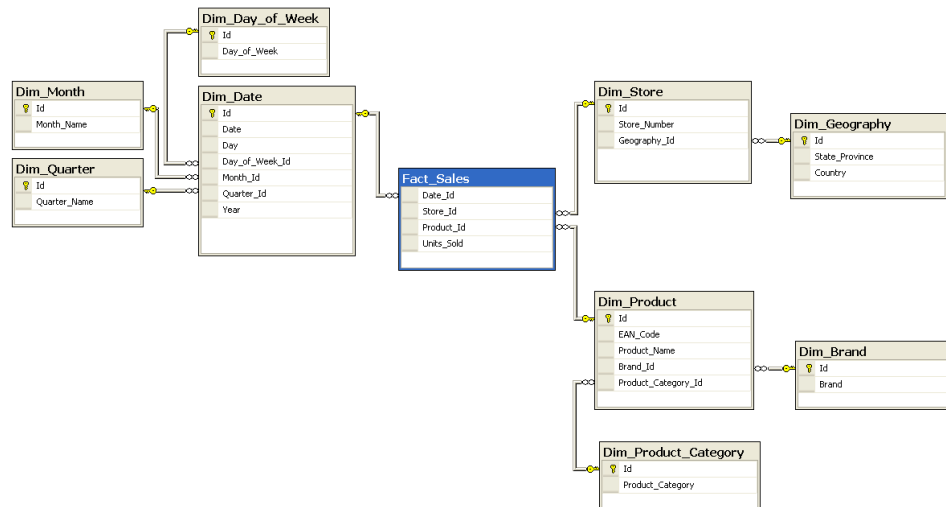
```
CREATE [UNIQUE|FULLTEXT|SPATIAL] INDEX index_name ON tbl (col);
```

- Foreign key
  - Index used to speed up JOIN queries
  - Not essential for JOIN queries

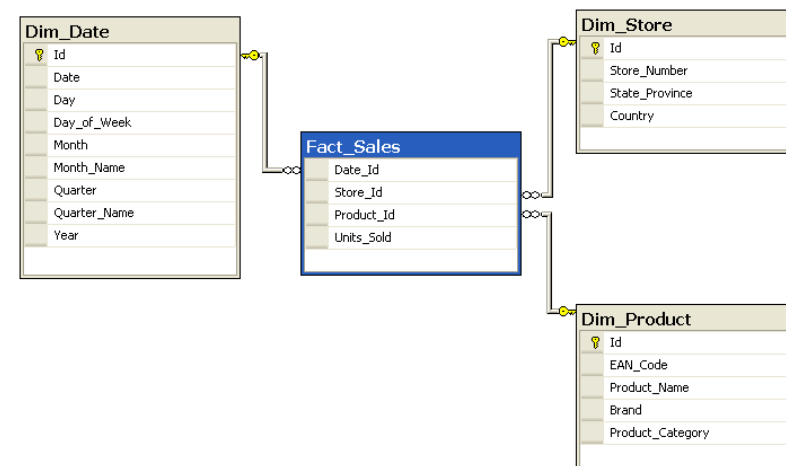
# Relational databases with MySQL

## Allow redundancy

- Schema with least redundancy isn't always the quickest
- Allow redundancy to reduce number of JOINS
- E.g. Data warehouse with real-time reporting
  - High efficiency is required
- Snowflake vs Star schema



Snowflake



Star

# Relational databases with MySQL

## SNOWFLAKE

- No redundancy
- Easy to maintain and change
- Complex queries
- Slower (more JOINS)
- Uses less space
- Bottom up

### SELECT

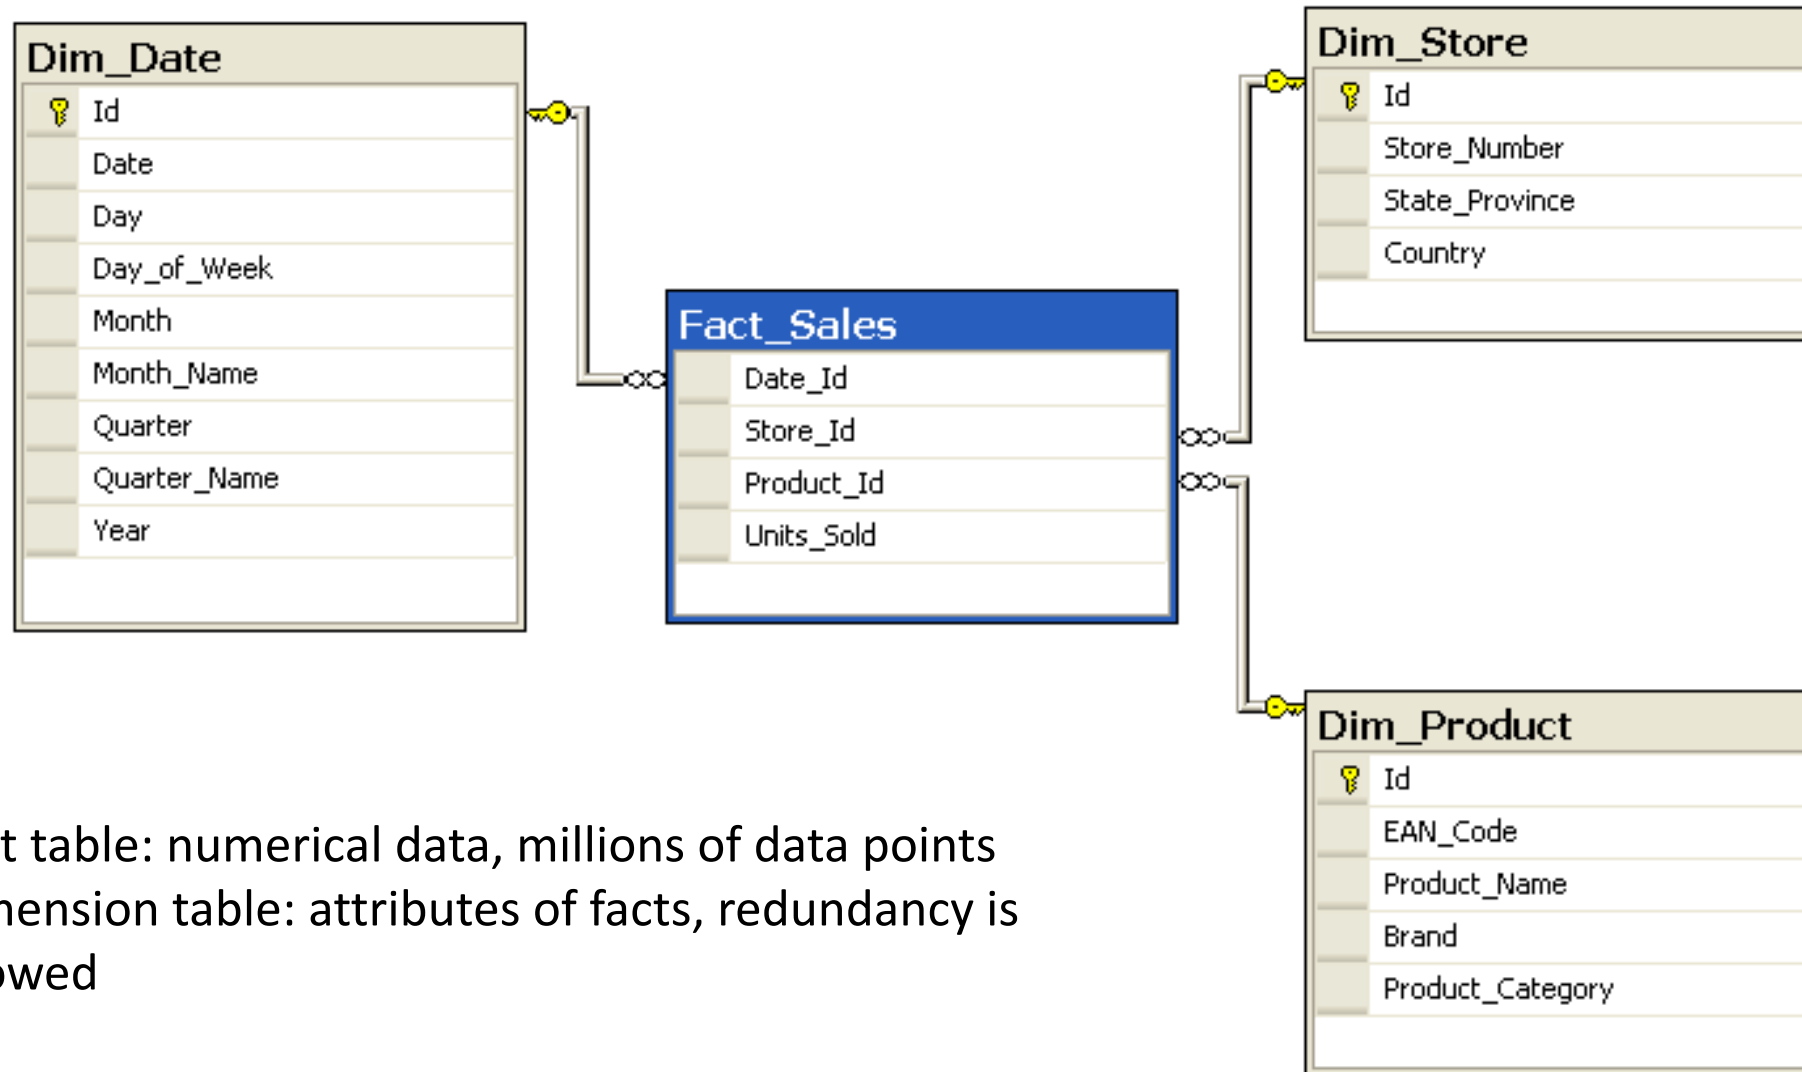
```
        B.Brand,  
        G.Country,  
        SUM(F.Units_Sold)  
FROM Fact_Sales F  
INNER JOIN Dim_Date D ON F.Date_Id = D.Id  
INNER JOIN Dim_Store S ON F.Store_Id = S.Id  
INNER JOIN Dim_Geography G ON S.Geography_Id = G.Id  
INNER JOIN Dim_Product P ON F.Product_Id = P.Id  
INNER JOIN Dim_Brand B ON P.Brand_Id = B.Id  
INNER JOIN Dim_Product_Category C ON  
P.Product_Category_Id = C.Id  
WHERE D.Year = 1997 AND C.Product_Category = 'tv'  
GROUP BY B.Brand, G.Country
```

## STAR

- Redundant data
- Less easy to maintain/change
- Lower query complexity
- Faster
- Uses more space (data is stored twice or more)
- Top down

```
SELECT    P.Brand,  
          S.Country AS Countries,  
          SUM(F.Units_Sold)  
FROM Fact_Sales F  
INNER JOIN Dim_Date D ON (F.Date_Id = D.Id)  
INNER JOIN Dim_Store S ON (F.Store_Id = S.Id)  
INNER JOIN Dim_Product P ON (F.Product_Id = P.Id)  
WHERE D.Year = 1997 AND P.Product_Category = 'tv'  
GROUP BY P.Brand, S.Country
```

# Relational databases with MySQL



- Fact table: numerical data, millions of data points
- Dimension table: attributes of facts, redundancy is allowed

# Relational databases with MySQL

## Database backup

- Dump a complete database into a tekst file  
`$ mysqldump [opt] db > db.sql`
- Includes
  - Statements for creating the database if the option `--databases` is used
  - Statements for creating tables, views, ...
  - Statements for inserting data
- Restore the database (may need to create first)  
`$ mysql db < db.sql`



# Relational databases with MySQL

## Exercises

- Create a separate dump file for each of your own databases
- Check the contents of each file
- For the bold ones
  - Drop your databases and recreate them using your dump files

# Relational databases with MySQL

## Rehearsal exercises

- Create a new database for your lab and include following data

- All trainings
  - Subject, duration
- All lab members
  - Name, lastname, birth\_date, training
- All equipment
  - Name, manufacturer, purchase\_date
- All experiments
  - Name, performed\_by, equipment\_used, date
- All results
  - Directory, experiment, status
- Fill with some data

```
CREATE TABLE Kit_order (  
    order_number CHAR(16) NOT NULL,  
    manufacturer VARCHAR(255),  
    kit_name VARCHAR(255),  
    supplier VARCHAR(255),  
    PRIMARY KEY (order_number),  
    FOREIGN KEY (manufacturer, kit_name)  
    REFERENCES Kit(manufacturer, name),  
    FOREIGN KEY (supplier)  
    REFERENCES Supplier(supplier_name));
```

```
CREATE TABLE Kit (  
    manufacturer VARCHAR(255) NOT NULL,  
    name VARCHAR(255) NOT NULL,  
    kit_cost DECIMAL(6,2),  
    buffer VARCHAR(255),  
    buffer_conc FLOAT,  
    enzyme VARCHAR(255),  
    enzyme_conc FLOAT,  
    nucl_mix VARCHAR(255),  
    nucl_conc FLOAT,  
    PRIMARY KEY (manufacturer, name));
```

# Relational databases with MySQL

## Rehearsal exercises (part 2)

- In your newly created database, search for
  - All experiment equipment purchased after 1st of January 1985
  - The number of experiments each lab member conducted
  - A list with all equipment used in a successful experiment
  - A list with all lab members that failed an experiment
  - Who followed wich trainings?
    - Number of participants per training